


Architecture Revisions



XScale is a trademark of Intel Corporation

Data Sizes and Instruction Sets

- The ARM is a 32-bit architecture.
 - When used in relation to the ARM:
 - **Byte** means 8 bits
 - **Halfword** means 16 bits (two bytes)
 - **Word** means 32 bits (four bytes)
 - Most ARM's implement two instruction sets
 - 32-bit ARM Instruction Set
 - 16-bit Thumb Instruction Set
 - Jazelle cores can also execute Java bytecode
-

Processor Modes

- The ARM has seven basic operating modes:
 - **User** : unprivileged mode under which most tasks run
 - **FIQ** : entered when a high priority (fast) interrupt is raised
 - **IRQ** : entered when a low priority (normal) interrupt is raised
 - **Supervisor** : entered on reset and when a Software Interrupt instruction is executed
 - **Abort** : used to handle memory access violations
 - **Undef** : used to handle undefined instructions
 - **System** : privileged mode using the same registers as user mode
-

The ARM Register Set

Current Visible Registers

Abort Mode

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)
r15 (pc)
cpsr
spsr

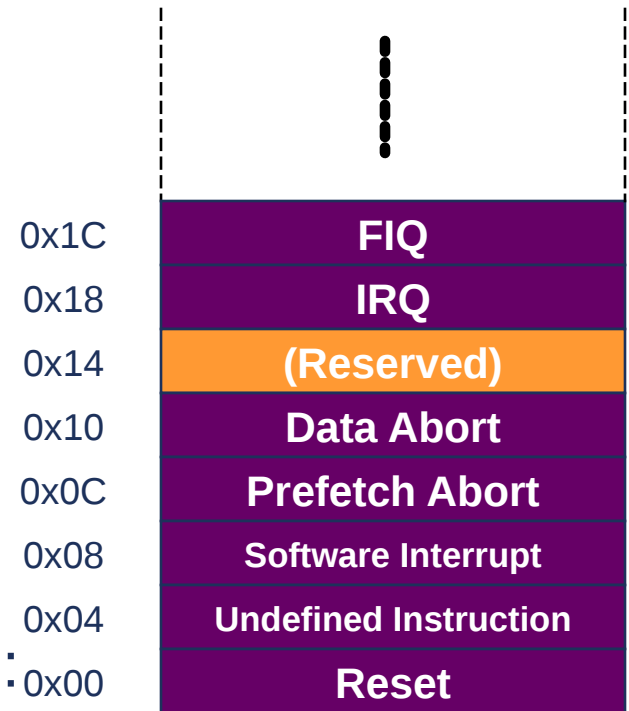
Banked out Registers

User	FIQ	IRQ	SVC	Undef
	r8			
	r9			
	r10			
	r11			
	r12			
r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)
r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)
	spsr	spsr	spsr	spsr

Exception Handling

- When an exception occurs, the ARM:
 - Copies CPSR into SPSR_<mode>
 - Sets appropriate CPSR bits
 - Change to ARM state
 - Change to exception mode
 - Disable interrupts (if appropriate)
 - Stores the return address in LR_<mode>
 - Sets PC to vector address
- To return, exception handler needs to:
 - Restore CPSR from SPSR_<mode>
 - Restore PC from LR_<mode>


This can only be done in ARM state.



Vector Table

Vector table can be at **0xFFFF0000** on ARM720T and on ARM9/10 family devices

Program Status Registers



■ Condition code flags

- N = **N**egative result from ALU
- Z = **Z**ero result from ALU
- C = ALU operation **C**arried out
- V = ALU operation **oV**erflowed

■ Sticky Overflow flag - Q flag

- Architecture 5TE/J only
- Indicates if saturation has occurred

■ J bit

- Architecture 5TEJ only
- J = 1: Processor in Jazelle state

■ Interrupt Disable bits.

- I = 1: Disables the IRQ.
- F = 1: Disables the FIQ.

■ T Bit

- Architecture xT only
- T = 0: Processor in ARM state
- T = 1: Processor in Thumb state

■ Mode bits

- Specify the processor mode


Program Counter (r15)

- **When the processor is executing in ARM state:**
 - All instructions are 32 bits wide
 - All instructions must be word aligned
 - Therefore the **pc** value is stored in bits [31:2] with bits [1:0] undefined (as instruction cannot be halfword or byte aligned)
 - **When the processor is executing in Thumb state:**
 - All instructions are 16 bits wide
 - All instructions must be halfword aligned
 - Therefore the **pc** value is stored in bits [31:1] with bit [0] undefined (as instruction cannot be byte aligned)
 - **When the processor is executing in Jazelle state:**
 - All instructions are 8 bits wide
 - Processor performs a word access to read 4 instructions at once
-

Conditional Execution and Flags

- ARM instructions can be made to execute conditionally by postfixing them with the appropriate condition code field.
 - This improves code density *and* performance by reducing the number of forward branch instructions.

```
CMP    r3, #0
BEQ    skip
ADD    r0, r1, r2
skip
```



```
CMP    r3, #0
ADDNE  r0, r1, r2
```

- By default, data processing instructions do not affect the condition code flags but the flags can be optionally set by using “S”. CMP does not need “S”.

loop

```
...
SUBS  r1, r1, #1
BNE  loop
```

decrement r1 and set flags

if Z flag clear then branch

Condition Codes

- The possible condition codes are listed below
 - Note AL is the default and does not need to be specified

Suffix	Description	Flags tested
EQ	Equal	Z=1
NE	Not equal	Z=0
CS/HS	Unsigned higher or same	C=1
CC/LO	Unsigned lower	C=0
MI	Minus	N=1
PL	Positive or Zero	N=0
VS	Overflow	V=1
VC	No overflow	V=0
HI	Unsigned higher	C=1 & Z=0
LS	Unsigned lower or same	C=0 or Z=1
GE	Greater or equal	N=V
LT	Less than	N!=V
GT	Greater than	Z=0 & N=V
LE	Less than or equal	Z=1 or N!=V
AL	Always	

Conditional execution examples

C source code

```
if (r0 == 0)
{
    r1 = r1 + 1;
}
else
{
    r2 = r2 + 1;
}
```

ARM instructions

unconditional

```
CMP r0, #0
BNE else
ADD r1, r1, #1
B end
else
    ADD r2, r2, #1
end
...
```

conditional

```
CMP r0, #0
ADDEQ r1, r1, #1
ADDNE r2, r2, #1
...
```

- 5 instructions
- 5 words
- 5 or 6 cycles

- 3 instructions
- 3 words
- 3 cycles


Data Processing Instructions

- Consist of :
 - Arithmetic: **ADD** **ADC** **SUB** **SBC** **RSB** **RSC**
 - Logical: **AND** **ORR** **EOR** **BIC**
 - Comparisons: **CMP** **CMN** **TST** **TEQ**
 - Data movement: **MOV** **MVN**
- These instructions only work on registers, NOT memory.
- Syntax:

<Operation>{<cond>}{S} Rd, Rn, Operand2

- Comparisons set flags only - they do not specify Rd
 - Data movement does not specify Rn
 - Second operand is sent to the ALU via barrel shifter.
-

Using a Barrel Shifter: The 2nd Operand



Register, optionally with shift operation

- Shift value can be either be:
 - 5 bit unsigned integer
 - Specified in bottom byte of another register.
- Used for multiplication by constant

Immediate value

- 8 bit number, with a range of 0-255.
 - Rotated right through even number of positions
- Allows increased range of 32-bit constants to be loaded directly into registers

Data Processing Exercise

1. How would you load the two's complement representation of -1 into Register 3 using one instruction?
 2. Implement an ABS (absolute value) function for a registered value using only two instructions.
 3. Multiply a number by 35, guaranteeing that it executes in 2 core clock cycles.
-

Data Processing Solutions

1. MOVN r6, #0

2. MOVS r7,r7 ; set the flags


RSBMI r7,r7,#0 ; if neg, r7=0-r7

3. ADD r9,r8,r8,LSL #2 ; r9=r8*5

RSB r10,r9,r9,LSL #3 ; r10=r9*7

Immediate constants

- No ARM instruction can contain a 32 bit immediate constant
 - All ARM instructions are fixed as 32 bits long
- The data processing instruction format has 12 bits available for operand2



Quick Quiz:
0xe3a004ff
MOV r0, #???

- 4 bit rotate value (0-15) is multiplied by two to give range 0-30 in steps of 2
- Rule to remember is
“8-bits rotated right by an even number of bit positions”

Loading 32 bit constants

- To allow larger constants to be loaded, the assembler offers a pseudo-instruction:
 - **LDR rd, =const**
- This will either:
 - Produce a **MOV** or **MVN** instruction to generate the value (if possible).or
 - Generate a **LDR** instruction with a PC-relative address to read the constant from a *literal pool* (Constant data area embedded in the code).
- For example
 - **LDR r0, =0xFF => MOV r0, #0xFF**
 - **LDR r0, =0x55555555 => LDR r0, [PC, #Imm12]**
 - ...
 - ...
 - DCD 0x55555555**
- This is the recommended way of loading constants into a register



Single register data transfer

LDR **STR** Word

LDRB **STRB** Byte

LDRH **STRH** Halfword

LDRSB Signed byte load

LDRSH Signed halfword load

- Memory system must support all access sizes
- Syntax:
 - **LDR**{<cond>}{<size>} Rd, <address>
 - **STR**{<cond>}{<size>} Rd, <address>

e.g. **LDREQB**

Address accessed

- Address accessed by LDR/STR is specified by a base register with an offset
- For word and unsigned byte accesses, offset can be:
 - An unsigned 12-bit immediate value (i.e. 0 - 4095 bytes)
LDR r0, [r1, #8]
 - A register, optionally shifted by an immediate value
LDR r0, [r1, r2]
LDR r0, [r1, r2, LSL#2]
- This can be either added or subtracted from the base register:
LDR r0, [r1, #-8]
LDR r0, [r1, -r2, LSL#2]
- For halfword and signed halfword / byte, offset can be:
 - An unsigned 8 bit immediate value (i.e. 0 - 255 bytes)
 - A register (unshifted)
- Choice of *pre-indexed* or *post-indexed* addressing
- Choice of whether to update the base pointer (pre-indexed only)
LDR r0, [r1, #-8]!

Load/Store Exercise

Assume an array of 25 words. A compiler associates y with $r1$. Assume that the base address for the array is located in $r2$. Translate this C statement/assignment using just three instructions:

`array[10] = array[5] + y;`

Load/Store Exercise Solution

`array[10] = array[5] + y;`

```
LDR    r3, [r2, #5]    ; r3 = array[5]
ADD     r3, r3, r1      ; r3 = array[5] + y
STR     r3, [r2, #10]   ; array[5] + y =
array[10]
```

Load and Store Multiples

- Syntax:

- `<LDM|STM>{<cond>}<addressing_mode> Rb{!}, <register list>`

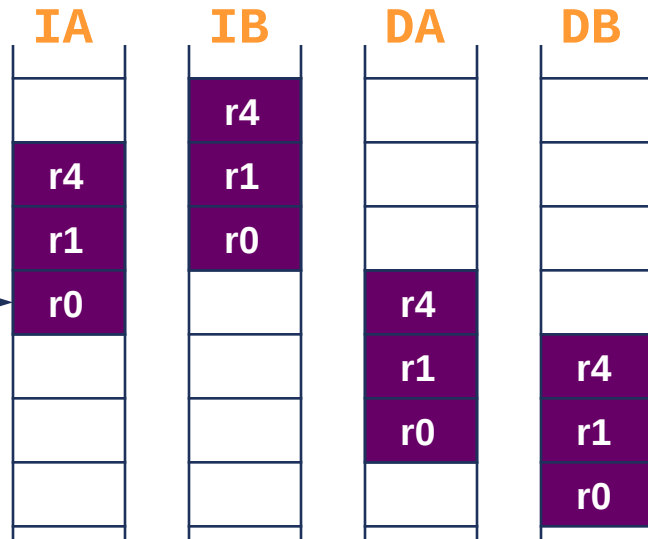
- 4 addressing modes:

- **LDMIA / STMIA** increment after
 - **LDMIB / STMIB** increment before
 - **LDMDA / STMDA** decrement after
 - **LDMDB / STMDB** decrement before

LDMxx r10, {r0,r1,r4}
STMxx r10, {r0,r1,r4}

Base Register (Rb)

r10



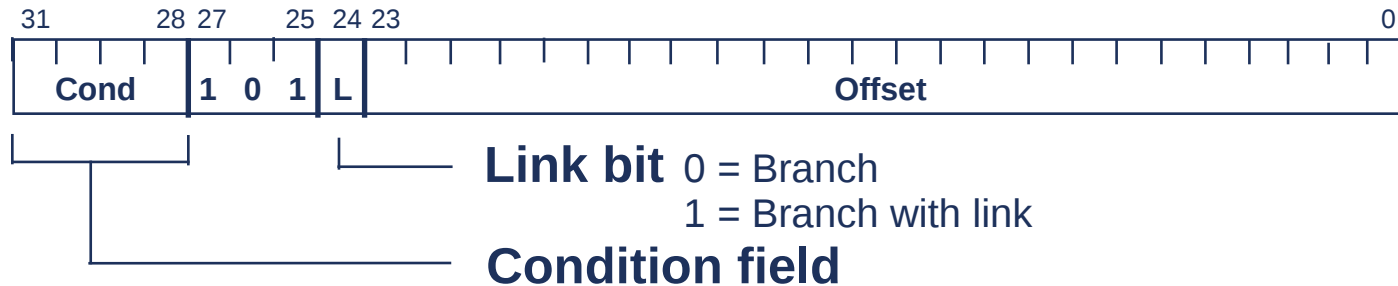
↑ Increasing Address

Multiply and Divide

- There are 2 classes of multiply - producing 32-bit and 64-bit results
- 32-bit versions on an ARM7TDMI will execute in 2 - 5 cycles
 - `MUL r0, r1, r2 ; r0 = r1 * r2`
 - `MLA r0, r1, r2, r3 ; r0 = (r1 * r2) + r3`
- 64-bit multiply instructions offer both signed and unsigned versions
 - For these instruction there are 2 destination registers
 - `[U|S]MULL r4, r5, r2, r3 ; r5:r4 = r2 * r3`
 - `[U|S]MLAL r4, r5, r2, r3 ; r5:r4 = (r2 * r3) + r5:r4`
- Most ARM cores do not offer integer divide instructions
 - Division operations will be performed by C library routines or inline shifts

Branch instructions

- Branch : `B{<cond>} label`
- Branch with Link : `BL{<cond>} subroutine_label`



- The processor core shifts the offset field left by 2 positions, sign-extends it and adds it to the PC
 - ± 32 Mbyte range
 - How to perform longer branches?

Register Usage

Arguments into function
Result(s) from function
otherwise corruptible
(Additional parameters
passed on stack)

Register

r0
r1
r2
r3

Register variables
Must be preserved

r4
r5
r6
r7
r8
r9/sb
r10/s1
r11

Scratch register
(corruptible)

r12

Stack Pointer
Link Register
Program Counter

r13/sp
r14/lr
r15/pc

The compiler has a set of rules known as a Procedure Call Standard that determine how to pass parameters to a function (see **AAPCS**)

CPSR flags may be corrupted by function call. Assembler code which links with compiled code must follow the AAPCS at external interfaces

The AAPCS is part of the new ABI for the ARM Architecture

- Stack base


- Stack limit if software stack checking selected

- SP should always be 8-byte (2 word) aligned


- R14 can be used as a temporary once value stacked

ARM Branches and Subroutines

- **B <label>**
 - PC relative. ± 32 Mbyte range.
- **BL <subroutine>**
 - Stores return address in LR
 - Returning implemented by restoring the PC from LR
 - For non-leaf functions, LR will have to be stacked



PSR access



- MRS and MSR allow contents of CPSR / SPSR to be transferred to / from a general purpose register or take an immediate value
 - MSR allows the whole status register, or just parts of it to be updated
- Interrupts can be enable/disabled and modes changed, by writing to the CPSR
 - Typically a read/modify/write strategy should be used:

```
MRS r0,CPSR      ; read CPSR into r0
BIC r0,r0,#0x80   ; clear bit 7 to enable IRQ
MSR CPSR_c,r0     ; write modified value to 'c' byte only
```

- In User Mode, all bits can be read but only the condition flags (_f) can be modified

Agenda

Introduction to ARM Ltd

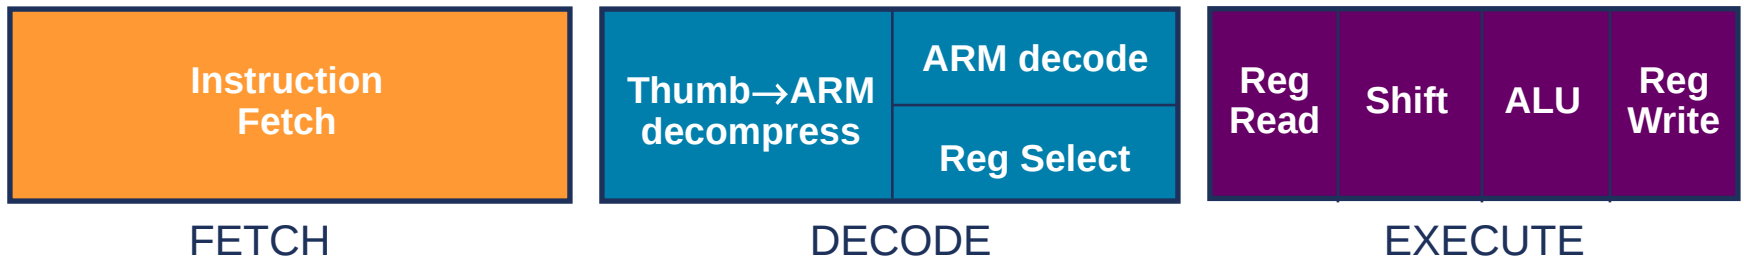
Fundamentals, Programmer's Model, and Instructions

- **Core Family Pipelines**

AMBA

Pipeline changes for ARM9TDMI

ARM7TDMI

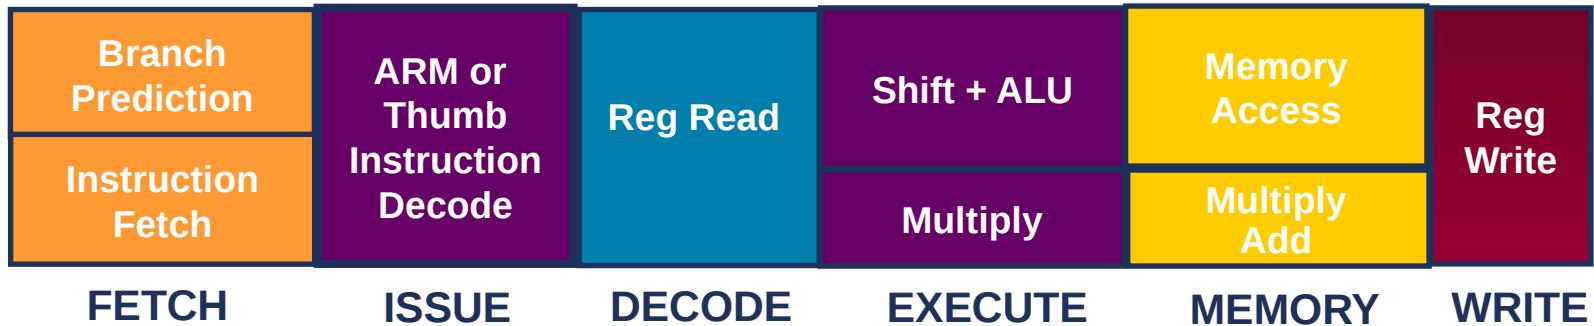


ARM9TDMI

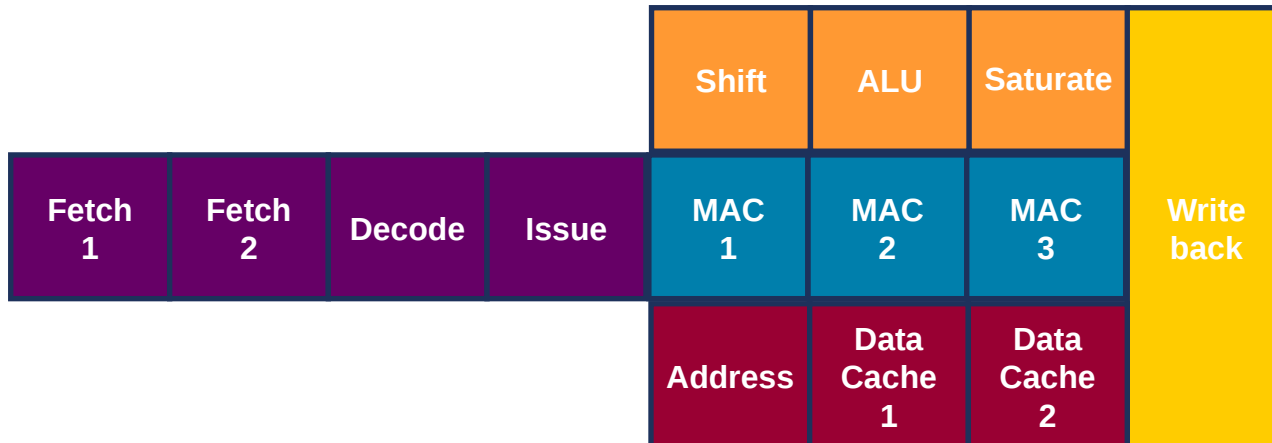


ARM10 vs. ARM11 Pipelines

ARM10



ARM11



Agenda

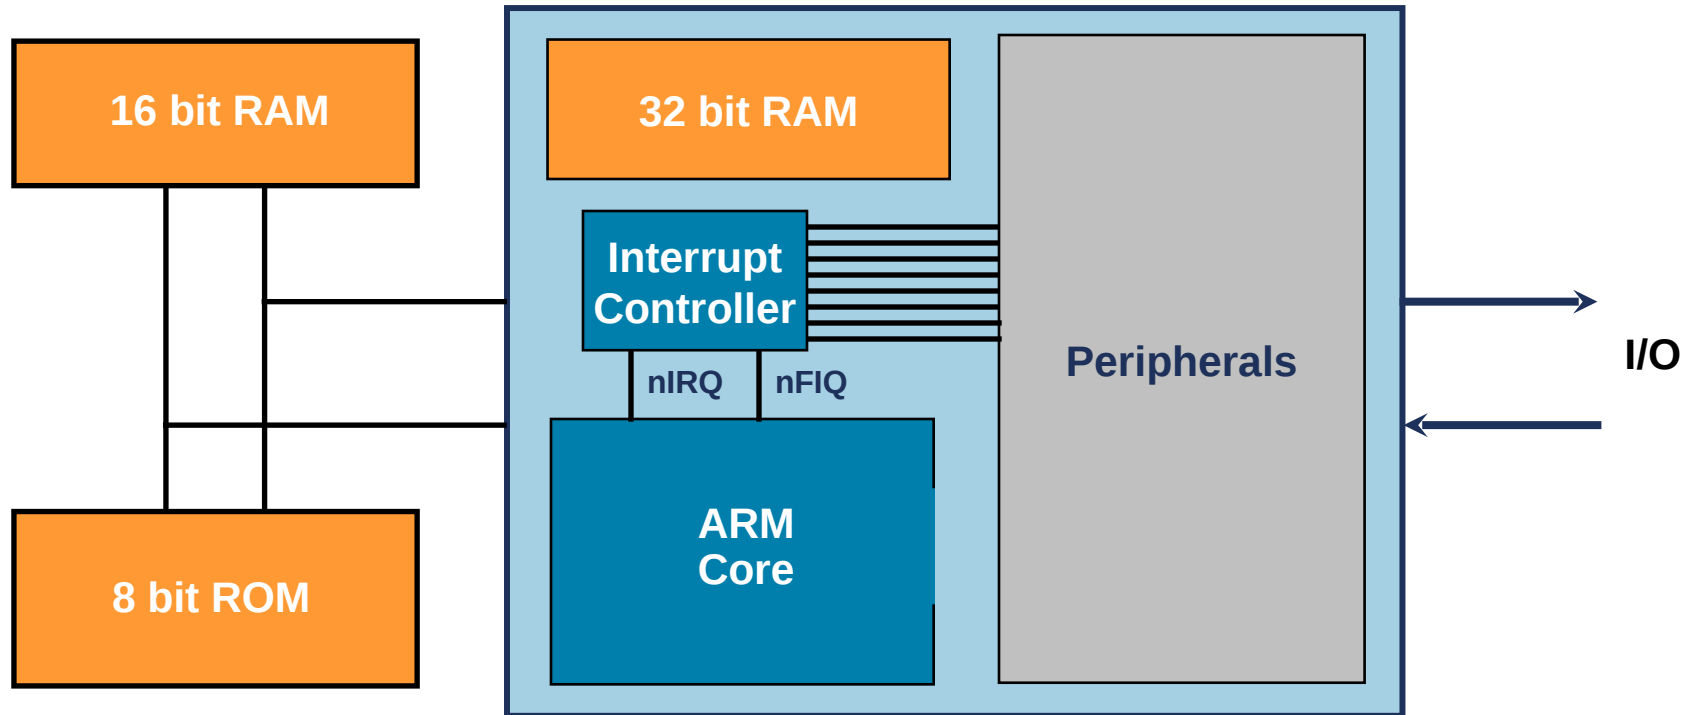
Introduction to ARM Ltd

Fundamentals, Programmer's Model, and Instructions

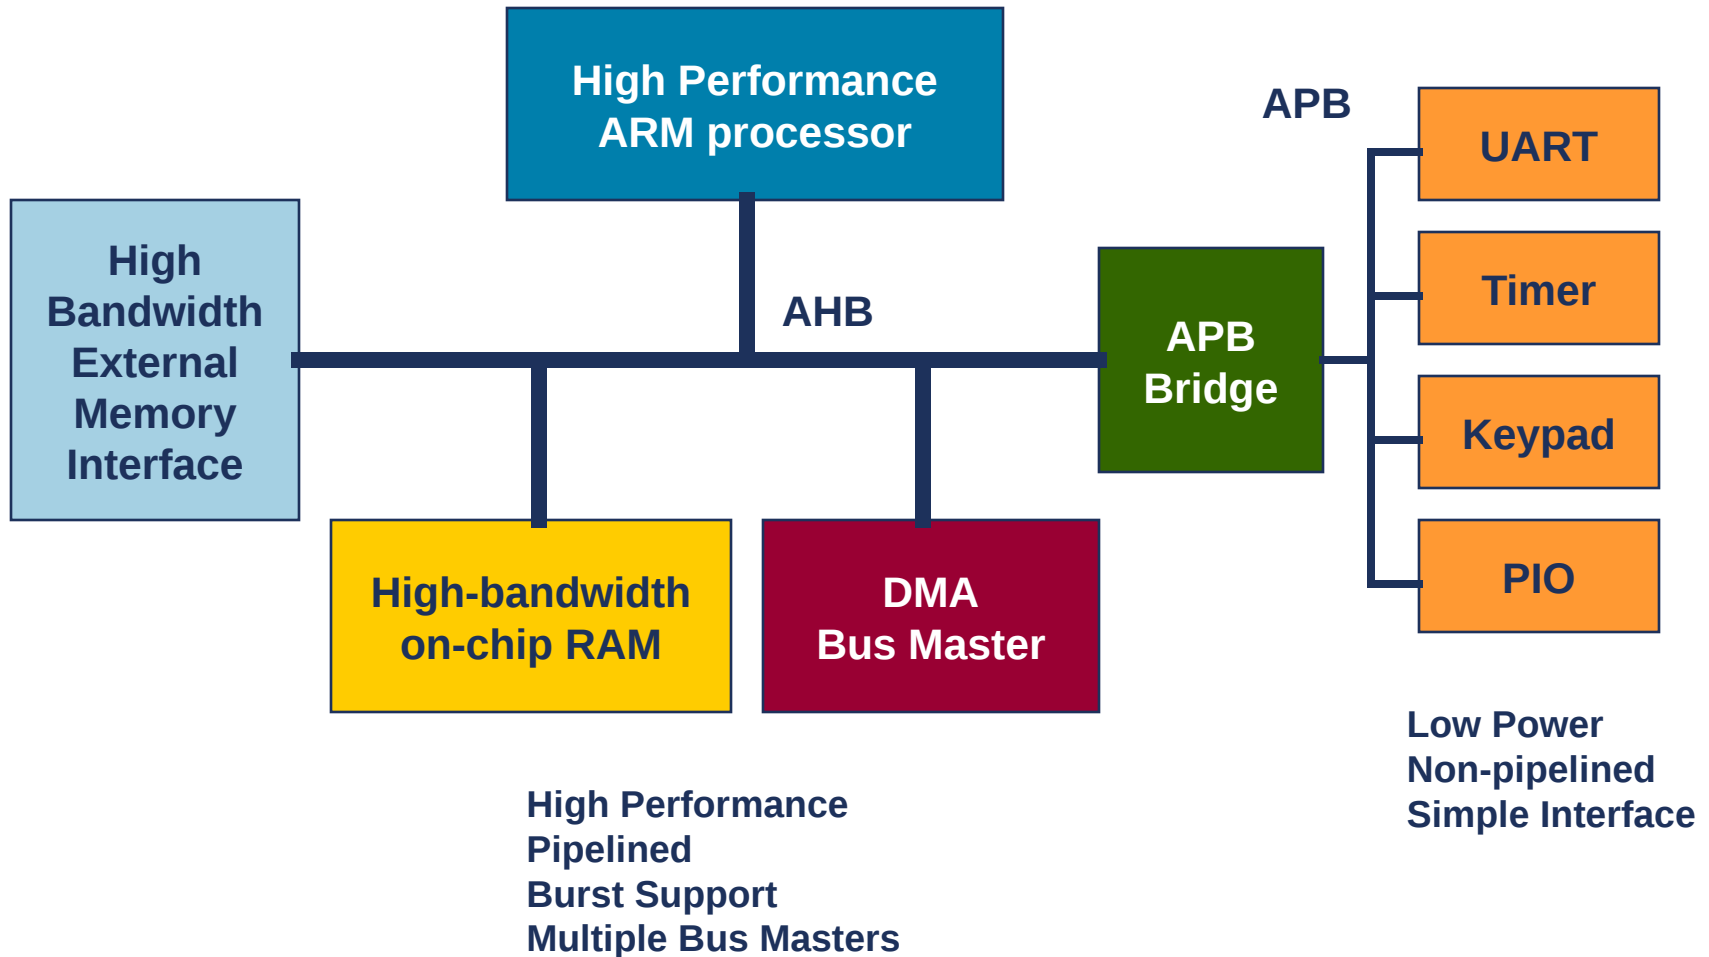
Core Family Pipelines

- **AMBA**

Example ARM-based System



An Example AMBA System



AHB Structure

